



5/8/2025.

L03 Spatial Filtering

In point operators




output point only depends on input point
not neighbours..


Introduction

- We use spatial filtering to enhance images, e.g., noise filtering.
- We can use the same technique to locate objects in images, called template matching. After completing this lesson, we would be able to process images using filters, as done in popular image processing software such as Photoshop.
- Intensity transformations affect the brightness (intensity level) of an image. The resulting value of a pixel after an intensity operation is just dependent on the original value of the same pixel.
- In contrast, the resulting value of a pixel after a spatial filtering operation depends on the neighborhood of the pixel in question. So, the space around the pixel in question matters, hence, the name spatial filtering.
- In linear spatial filtering we use the 2-D convolution operation.

Spatial Filtering



A 3×3 window (kernel or filter)



Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	18	90	90	18	90	0	
3	0	0	18	90	18	18	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

kernel

	0	1	2	3	4	5	6	7	j
0									
1									
2									
3									
4									
5									
i									


Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									


1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1		20	30						
2									
3									
4									
5									
i									

Spatial Filtering



1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9



Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1		20	30	40	40				
2									
3									
4									
5									
i									

Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1		20	30	40	40	40			
2									
3									
4									
5									
i									

Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1		0	30	40	40	40	30		
2									
3									
4									
5									
i									

Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9


	0	1	2	3	4	5	6	7	j
0									
1		0	30	40	40	40	30		
2		40	60	90	90	80	50		
3		50	90	90	90	80	50		
4		20	30	50	50	40	20		
5									
i									

when we try to fill the other cells
the kernel will go outside
So we can use padding

Practical Details: Dealing with Image Boundaries

To control the size of the output, we need to use *padding*

for a 5×5 kernel the border will be 2 cells wide



Source: Lazebnik with a slight modification.

x	x	x
x	x	x
x	x	x

Kernel

x	x	x
x	x	x
x	x	x

we essentially
multiply
and
add

This is the dot
product

$$\begin{bmatrix} \underline{h} \\ x \\ x \\ x \\ \vdots \\ x \end{bmatrix} \cdot \begin{bmatrix} \underline{f} \\ x \\ x \\ x \\ \vdots \\ x \end{bmatrix}$$

$$\underline{h} \cdot \underline{f} = |h| |f| \cos \theta$$


max when $\cos \theta = 1$
in other words the kernel
matches the image.

So the output is highest when the image
looks like the kernel
kernel looks for similar areas

Practical details: Dealing with Image Boundaries

- To control the size of the output, we need to use *padding*
- What values should we pad the image with?
 - Zero pad (or clip filter) *BORDER_CONSTANT*
 - Wrap around *BORDER_WRAP*
 - Copy edge *BORDER_REPLICATE*
 - Reflect across edge *BORDER_REFLECT_101*

Result
[0, 0, 1, 2, 3, 4, 0, 0]
[3, 4, 1, 2, 3, 4, 1, 2]
[1, 1, 1, 2, 3, 4, 4, 4]
[3, 2, 1, 2, 3, 4, 3, 2]




Examples of Effect of Kernel Choices




Original



Averaging



Sobel vertical



Sobel horizontal

In sobel implementation we use correlation rather than convolution.

Forward difference?

Dark

Bright

forward

backward

$$\begin{bmatrix} -1 \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$$

we do

$$\begin{bmatrix} -1 & 1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

Then we emphasize the middle since we are finding the middle pixel.

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

Averaging

dark

-1	-2	-1
0	0	0
1	2	1

Sobel vertical

-1	0	1
-2	0	2
-1	0	1

Sobel horizontal

bright

Averaging Using cv.filter2D

```
import cv2 as cv
import numpy as np
from matplotlib import pyplot as plt
```

Load a grayscale and reduce size by half on both dimensions
we: quickly load a smaller grayscale version of an image

```
im = cv.imread('images/sigiriya.jpg', cv.IMREAD_REDUCED_GRAYSCALE_2)
assert im is not None
```

```
kernel = np.ones((3,3),np.float32)/9
imavg = cv.filter2D(im, cv.CV_32F, kernel)
```

```
fig, axes = plt.subplots(1,2, sharex='all', sharey='all', figsize=(18,9))
axes[0].imshow(im, cmap='gray')
axes[0].set_title('Original')
axes[0].set_xticks([]), axes[0].set_yticks([])
axes[1].imshow(imavg, cmap='gray')
axes[1].set_title('Averaging')
axes[1].set_xticks([]), axes[1].set_yticks([])
plt.show()
```

In filter2D it gets border reflect as padding
and it does not flip the kernel

Sobel Filtering Using cv.filter2D

```
import cv2 as cv
import numpy as np
from matplotlib import pyplot as plt

im = cv.imread('images/sigiriya.jpg', cv.IMREAD_REDUCED_GRAYSCALE_2)
assert im is not None

sobel_x = np.array([[ -1,  -2,  -1], [ 0,  0,  0], [ 1,  2,  1]])
sobel_y = np.array([[ -1,  0,  1], [-2,  0,  2], [-1,  0,  1]])

im_x = cv.filter2D(im, cv.CV_64F, sobel_x)
im_y = cv.filter2D(im, cv.CV_64F, sobel_y)

fig, ax = plt.subplots(1,2, sharex='all', sharey='all', figsize=(18,9))
ax[0].imshow(im_x, cmap='gray')
ax[0].set_title('Sobel X')
ax[0].set_xticks([]), ax[0].set_yticks([])
ax[1].imshow(im_y, cmap='gray')
ax[1].set_title('Sobel Y')
ax[1].set_xticks([]), ax[1].set_yticks([])
plt.show()
```

vertical

horizontal

Central difference explained

In 1-D central difference

$$f'(x) = \frac{f(x+1) - f(x-1)}{2}$$

in sobel we use $\begin{bmatrix} -1 & 0 & +1 \end{bmatrix}$
it is similar

in the image we $\oplus$ the $x+1$ value and
 $\ominus$ the $x-1$ value.

also in the other direction of the
sobel we have $\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix}$ which gives

a smoothing effect on the vertical axis.

even with just $\begin{bmatrix} -1 & 0 & +1 \end{bmatrix}$ we will
get the sobel effect but it will
be much more noisy.

for sobel we can use zero padding
- the math derivatives are 0
outside the boundary
- to prevent artificial gradients
- black borders not appeared
area.

4. When to Use Other Modes

Padding Mode	Ideal Use Case	Sobel Suitability
<code>constant</code> (zero)	Edge detection, technical imaging	★★★★★
<code>edge</code> (replicate)	Natural photo processing	★★★★☆
<code>reflect</code>	Texture analysis, seamless processing	★★★★☆
<code>wrap</code>	Periodic signals (like FFT)	★★★★☆

Convolution and Correlation

- Spatial filtering is, in fact, convolution.
- As the filters are typically symmetric—i.e., a 180° rotation results in the same kernel—correlation is equivalent to convolution.
- Correlation is also the scalar product between the kernel and the underlying image patch. Therefore, it is a measure of similarity between the kernel and the underlying image patch.
- As a result, when the kernel and the patch are “similar”, the output is high. In view of this, **spatial filtering seeks for patches in the image that are similar to the kernel.**
- Implementing filtering using loops (four nested for loops) in a non-C fashion is inefficient. Instead, use filter2D.

we did not flip the kernel. why.
convolution and correlation is similar for symmetric
kernels.

Convolution and Correlation

The convolution sum expression that we learned in signal and systems is

$$y[n] = x[n] * b[n] = \sum_{k=-\infty}^{\infty} x[k]b[n-k].$$

In 2-D, as applicable in image processing, correlation sum with a kernel $w[m, n]$ with non-zero values in $(m, n) \in ([-a, a], [-b, b])$ is

$$(w * f)[m, n] = w[m, n] * f[m, n] = \sum_{s=-a}^a \sum_{t=-b}^b w[s, t]f[m-s, n-t].$$

Correlation:

$$(w \circledast f)[m, n] = w[m, n] \circledast f[m, n] = \sum_{s=-a}^a \sum_{t=-b}^b w[s, t]f[m+s, n+t].$$

} looks similar

Example

Consider the image

$$f = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

and the filtering kernel

$$w = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}$$

. Appropriately pad the image. Carry out a. correlation b. convolution.

after flipping or not, f remains the same.

pad it to be 7×7

$f =$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

c) $\otimes$

8 8 7

6 5 4

3 2 1

Example

Consider the image

$$f_{\text{padded}} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \quad w = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}$$

Correlation result and convolution result, respectively

$$(w \circledast f) = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 9 & 8 & 7 & 0 \\ 0 & 6 & 5 & 4 & 0 \\ 0 & 3 & 2 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \quad (w * f) = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 3 & 0 \\ 0 & 4 & 5 & 6 & 0 \\ 0 & 7 & 8 & 9 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

if we do not flip, the
result will be flipped

so we need to
flip first to get it correct.

Convolution: Key Properties

1. Linearity: $\text{filter}(f_1 + f_2) = \text{filter}(f_1) + \text{filter}(f_2)$
2. Shift invariance: same behavior regardless of pixel location:
 $\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$
3. Theoretical result: any linear shift-invariant operator can be represented as a convolution.

Other Properties

1. Commutative: $a * b = b * a$
2. Conceptually no difference between filter and signal
3. Associative: $a * (b * c) = (a * b) * c$
4. Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
5. This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
6. Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
7. Scalars factor out: $ka * b = a * kb = k(a * b)$
8. Identity: unit impulse $e = [\dots, 0, 0, 1, 0, 0, \dots]$, $a * e = a$

Suppose we have a two 3×3 kernels
and a 1000×1000 image

$$\text{Total computations} = 3^2 \times 1000^2 \times 2$$

but due to associativity

$$\underbrace{a * (b * c)}_{\substack{3 \times 3 \quad 3 \times 3 \quad 1000^2}} = (a * b) * c$$

$$3^2 \times 3^2 + 3^2 \times 1000^2$$

At the Edge

The filter window falls off the edge of the image.

Need to extrapolate the image.

Methods: various border types, image boundaries are denoted with '|'

- BORDER_REPLICATE: aaaaaa | abcdefgh | hhhhhhhh
- BORDER_REFLECT: fedcba | abcdefgh | hgfedcb
- BORDER_REFLECT_101: gfedcb | abcdefgh | gfedcba
- BORDER_WRAP: cdefgh | abcdefgh | abcdefg
- BORDER_CONSTANT: iiiiii | abcdefgh | iiiiii with some specified 'i'

Box Filter vs. Gaussian Filter

- What's wrong with this filtering operation?
- What's the solution?



(a) Original



(b) Kernel (much zoomed)



(c) Box Filtered

Gaussian has better quality
• box has sharp cutoff and in freq domain it has sinc response which has strong side lobes (Gibbs phenomenon)
• gaussian is spatially aware / smooth weight fall off / center contribute more


Box Filter vs. Gaussian Filter



(a) Original



(b) Box Filtered






(c) Gaussian Filtered

Gaussian Kernel

$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

one std (pointing to σ)

mean = 0



- Constant factor at front makes volume sum to 1 (can be ignored when computing the filter values, as we should renormalize weights to sum to 1 in any case).
- The Gaussian function has infinite support, but discrete filters use finite kernels.

Compute a 2D 3×3 gaussian by hand.
(un-normalized)

$$g(x, y) = e^{-\frac{(x^2 + y^2)}{2\sigma^2}} \quad \sigma = 1$$

$$= e^{-\frac{x^2 + y^2}{2}}$$

compute for a 3×3 grid

$(-1, 1)$ $\frac{1}{e}$	$(0, 1)$ $\frac{1}{\sqrt{e}}$	$\frac{1}{e}$
$\frac{1}{\sqrt{e}}$	$(0, 0)$ 1	$\frac{1}{\sqrt{e}}$
$\frac{1}{e}$	$\frac{1}{\sqrt{e}}$	$\frac{1}{e}$

$$\frac{1}{\sqrt{e}} = 0.607$$

$$\frac{1}{e} = 0.368$$

Now we have to normalize it by dividing by the sum of the kernel.


Otherwise an all white area can overflow.

Choosing Gaussian Kernel Width

$$g(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

low
blur

Rule of thumb: set the filter half-width to about 3σ .




for $\sigma = 3 \Rightarrow 3\sigma = 9$
filter = $(9 \times 2) + 1 = 19$

doesn't go to 0

Separability of the Gaussian Filter

$$G(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right) \cdot \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{y^2}{2\sigma^2}\right)$$



=



*



Note: convolution and matrix multiplication (outer product) are the same in this case

$$(1 \times 1) \times 2$$

11x11 ⇒ complexity

- The 2-D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y .
- In this case, the two Gaussians are the (identical) 1-D Gaussians.
- Separability means that a 2-D convolution can be reduced to two 1-D convolutions (one among rows and one among columns).
- What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?
- What is the complexity if the kernel is separable?

$$O(n^2 m^2)$$

$$m^2 n^2$$


$$O(n^2 m)$$

$$m n^2 + m n^2$$

image / kernel

Gaussian Noise


- Mathematical model: sum of many independent factors
- **Good for small standard deviations**
- Assumption: independent, zero-mean noise



$$f(x, y) = \overbrace{f(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$


when σ increase, noise level increase. more dots.
 when the smoothing kernel size increase it reduces
 noise effect but blurs the image more.



Reducing Gaussian Noise
 Smoothing with larger standard deviations
 suppresses noise but also blurs the image.


Reducing Salt-and-Pepper Noise: Median Filtering

not linear



Inability to reduce salt and pepper noise with Gaussian filtering

*bigger window
→ remove more noise but
more detail loss*



A median filter operates over a window by selecting the median intensity in the window.

Is median filtering linear?

NO

Source: K. Grauman

50 50 200
 50 50 200 60 70 80
 45 45 65 75 5 85
 80 90 100

5 60 70 70 (65) 60 65 60 60

50 50 50 (65) 65 65 200 200

Bilateral Filter

Blur only a


Mathematical Formulation

$$I_{\text{filtered}}(x) = \frac{1}{W_p} \sum_{x_i \in \Omega} G_{\sigma_s}(\underbrace{\|x_i - x\|}_{\substack{\text{close say 1} \\ \text{completely blur}}}) \cdot G_{\sigma_r}(\underbrace{\|I(x_i) - I(x)\|}_{\substack{\text{what we did} \\ \text{distance between pixels.} \\ 150 \Rightarrow e^{-150} \approx 0 \text{ no blur.}}})$$

Original



Gaussian Filtered



Bilateral Filtered

